

# Exploratory Data Analysis of StackOverFlow Survery on Salaries of Various Coding Background Professions

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

import altair as alt
from sklearn.model_selection import train_test_split

alt.data_transformers.enable('data_server')
alt.renderers.enable('mimetype')
```

```
Out[1]: RendererRegistry.enable('mimetype')
```

## Functions used

```
In [2]: # Returns True if 'data' in a string , else returns False
def data_in_text(x):
    if type(x) is str:
        return 'data' in x
    else:
        return False

# Returns float values for different string inputs
def convert2float(x):
    if x == 'More than 50 years' :
        return float(50)
    elif x == 'Less than 1 year':
        return float(0)
    else:
        return float(x)
```

## Summary of the dataset

The goal of the project is to create a machine learning model by using the job related information from the collected dataset, to predict the target `ConvertedCompYearly`, which is the converted yearly compensation for the Data scientist/analysis related jobs in the US and Canada.

The final model will be a regression model as the target will be a continuous variable.

```
In [3]: survey_df = pd.read_csv("data/survey_results_public.csv")
```

```
In [4]: survey_df.info()
```

[illegible]

0	ResponseId	73268	non-null	int64
1	MainBranch	73268	non-null	object
2	Employment	71709	non-null	object
3	RemoteWork	58958	non-null	object
4	CodingActivities	58899	non-null	object
5	EdLevel	71571	non-null	object
6	LearnCode	71580	non-null	object
7	LearnCodeOnline	50685	non-null	object
8	LearnCodeCoursesCert	29389	non-null	object
9	YearsCode	71331	non-null	object
10	YearsCodePro	51833	non-null	object
11	DevType	61302	non-null	object
12	OrgSize	51039	non-null	object
13	PurchaseInfluence	50969	non-null	object
14	BuyNewTool	67963	non-null	object
15	Country	71771	non-null	object
16	Currency	51264	non-null	object
17	CompTotal	38422	non-null	float64
18	CompFreq	44425	non-null	object
19	LanguageHaveWorkedWith	70975	non-null	object
20	LanguageWantToWorkWith	67027	non-null	object
21	DatabaseHaveWorkedWith	60121	non-null	object
22	DatabaseWantToWorkWith	51014	non-null	object
23	PlatformHaveWorkedWith	49924	non-null	object
24	PlatformWantToWorkWith	40415	non-null	object
25	WebframeHaveWorkedWith	53544	non-null	object
26	WebframeWantToWorkWith	46122	non-null	object
27	MiscTechHaveWorkedWith	44992	non-null	object
28	MiscTechWantToWorkWith	36810	non-null	object
29	ToolsTechHaveWorkedWith	54171	non-null	object
30	ToolsTechWantToWorkWith	46566	non-null	object
31	NEWCollabToolsHaveWorkedWith	70347	non-null	object
32	NEWCollabToolsWantToWorkWith	64108	non-null	object
33	OpSysProfessional use	65503	non-null	object
34	OpSysPersonal use	70963	non-null	object
35	VersionControlSystem	71379	non-null	object
36	VCInteraction	68156	non-null	object
37	VCHostingPersonal use	0	non-null	float64
38	VCHostingProfessional use	0	non-null	float64
39	OfficeStackAsyncHaveWorkedWith	46223	non-null	object
40	OfficeStackAsyncWantToWorkWith	32072	non-null	object
41	OfficeStackSyncHaveWorkedWith	62128	non-null	object
42	OfficeStackSyncWantToWorkWith	47688	non-null	object
43	Blockchain	71071	non-null	object
44	NEWSOSites	71365	non-null	object
45	SOVisitFreq	70961	non-null	object
46	SOAccount	71572	non-null	object
47	SOPartFreq	58229	non-null	object
48	SOComm	71408	non-null	object
49	Age	70946	non-null	object
50	Gender	70853	non-null	object
51	Trans	70315	non-null	object
52	Sexuality	66565	non-null	object
53	Ethnicity	69474	non-null	object
54	Accessibility	67244	non-null	object
55	MentalHealth	66447	non-null	object
56	TBranch	52670	non-null	object
57	ICorPM	36283	non-null	object
58	WorkExp	36769	non-null	float64
59	Knowledge_1	35804	non-null	object
60	Knowledge_2	34973	non-null	object
61	Knowledge_3	35133	non-null	object
62	Knowledge_4	35097	non-null	object
63	Knowledge_5	35014	non-null	object
64	Knowledge_6	34991	non-null	object
65	Knowledge_7	34977	non-null	object

```

66 Frequency_1          35371 non-null object
67 Frequency_2          35344 non-null object
68 Frequency_3          34515 non-null object
69 TimeSearching        36198 non-null object
70 TimeAnswering        36022 non-null object
71 Onboarding           35679 non-null object
72 ProfessionalTech     34906 non-null object
73 TrueFalse_1          35819 non-null object
74 TrueFalse_2          35715 non-null object
75 TrueFalse_3          35749 non-null object
76 SurveyLength         70444 non-null object
77 SurveyEase           70508 non-null object
78 ConvertedCompYearly  38071 non-null float64
dtypes: float64(5), int64(1), object(73)
memory usage: 44.2+ MB

```

As we can see there are 79 columns and 73268 rows.

```
In [5]: survey_df['Country'].value_counts()[0:10].to_frame().style.set_caption('Table 1. Counts o
```

```
Out[5]: Table 1. Counts of observation for top ten countries
```

	Country
	<b>United States of America</b> 13543
	<b>India</b> 6639
	<b>Germany</b> 5395
	<b>United Kingdom of Great Britain and Northern Ireland</b> 4190
	<b>Canada</b> 2490
	<b>France</b> 2328
	<b>Brazil</b> 2109
	<b>Poland</b> 1732
	<b>Netherlands</b> 1555
	<b>Spain</b> 1521

Since we want to focus on the data points for US and Canada, data scientist/analysis related jobs, we filter the dataset by `United States of America` and `Canada`.

```
In [6]: north_america_data = survey_df.query("Country == 'United States of America' or Country =
```

```
In [7]: north_america_data.shape
```

```
Out[7]: (16033, 79)
```

Since we don't need all the features from the dataset, thus we will only keep the remaining ones for the prediction purpose:

- MainBranch
- Employment
- RemoteWork
- EdLevel
- YearCode
- YearCodePro
- DevType

- OrgSize
- Country
- LanguageHaveWorkedWith
- DatabaseHaveWorkedWith
- PlatformHaveWorkedWith
- WebframeHaveWorkedWith
- MiscTechHaveWorkedWith
- ToolTechHaveWorkedWith
- NEWCollabToolsHaveWorkedWith
- OpSysProfessional use
- VersionControlSystem
- VCInteraction
- OfficeStackAsyncHaveWorkedWith
- Age
- WorkExp
- Icorp

```
In [8]: cols_to_choose = ['MainBranch',
    'Employment',
    'RemoteWork',
    'EdLevel',
    'YearsCode',
    'YearsCodePro',
    'DevType',
    'OrgSize',
    'Country',
    'LanguageHaveWorkedWith',
    'DatabaseHaveWorkedWith',
    'PlatformHaveWorkedWith',
    'WebframeHaveWorkedWith',
    'MiscTechHaveWorkedWith',
    'ToolsTechHaveWorkedWith',
    'NEWCollabToolsHaveWorkedWith',
    'OpSysProfessional use',
    'VersionControlSystem',
    'VCInteraction',
    'OfficeStackAsyncHaveWorkedWith',
    'Age',
    'WorkExp',
    'ICorPM',
    'ConvertedCompYearly']
```

```
In [9]: north_america_data = north_america_data[cols_to_choose]
```

```
In [10]: multianswerq_cols = [
    'DevType',
    'LanguageHaveWorkedWith',
    'DatabaseHaveWorkedWith',
    'PlatformHaveWorkedWith',
    'WebframeHaveWorkedWith',
    'MiscTechHaveWorkedWith',
    'ToolsTechHaveWorkedWith',
    'NEWCollabToolsHaveWorkedWith',
    'OpSysProfessional use',
    'VCInteraction',
    'VersionControlSystem',
    'OfficeStackAsyncHaveWorkedWith',
    'Employment']
```

# Split data set into training and test splits

---

```
In [11]: train_df, test_df = train_test_split(north_america_data, test_size=0.2, random_state=123)

train_df = train_df.dropna(subset=['ConvertedCompYearly'])
```

## Data Wrangling

---

```
In [12]: train_df['YearsCode'] = train_df['YearsCode'].apply(lambda x: convert2float(x))

train_df['YearsCodePro'] = train_df['YearsCodePro'].apply(lambda x: convert2float(x))
```

## EDA on the Training Dataset

---

### Coding for all the charts

```
In [13]: converted_comp_hist = alt.Chart(train_df).mark_bar().encode(
    x=alt.X('ConvertedCompYearly', bin=alt.Bin(maxbins=30)),
    y='count()'
)

years_code = alt.Chart(train_df).mark_rect().encode(
    alt.X('YearsCode', bin=alt.Bin(maxbins=40)),
    alt.Y('ConvertedCompYearly', bin=alt.Bin(maxbins=40)),
    alt.Color('count()')
)

years_code_pro = alt.Chart(train_df).mark_rect().encode(
    alt.X('YearsCodePro', bin=alt.Bin(maxbins=40)),
    alt.Y('ConvertedCompYearly', bin=alt.Bin(maxbins=40)),
    alt.Color('count()')
)

work_exp = alt.Chart(train_df).mark_rect().encode(
    alt.X('WorkExp', bin=alt.Bin(maxbins=40)),
    alt.Y('ConvertedCompYearly', bin=alt.Bin(maxbins=40)),
    alt.Color('count()')
)


columns = ['Country', 'ICorPM', 'RemoteWork', 'OrgSize', 'EdLevel']


multi_bar_plot = alt.Chart(train_df).mark_bar().encode(
    y=alt.Y(alt.repeat(), sort='x'),
    x='median(ConvertedCompYearly)'
).repeat(columns, columns=2)

corr_table = train_df.corr().style.background_gradient().set_caption('Table 2. Correlati
```


```
In [14]: # only for plotting
plotting_dataset = train_df.copy()
for col in multianswerq_cols:
    plotting_dataset = plotting_dataset.dropna(subset=[col])
```

```
plotting_dataset[col] = plotting_dataset[col].apply(lambda x: x.split(';'))
boom_data = plotting_dataset.explode(col)
boom_data[col].value_counts().plot.barh(title = col, figsize = [15, 15])
plt.show()
```







DatabaseHaveWorkedWith







PlatformHaveWorkedWith










ToolsTechHaveWorkedWith










VCInteraction





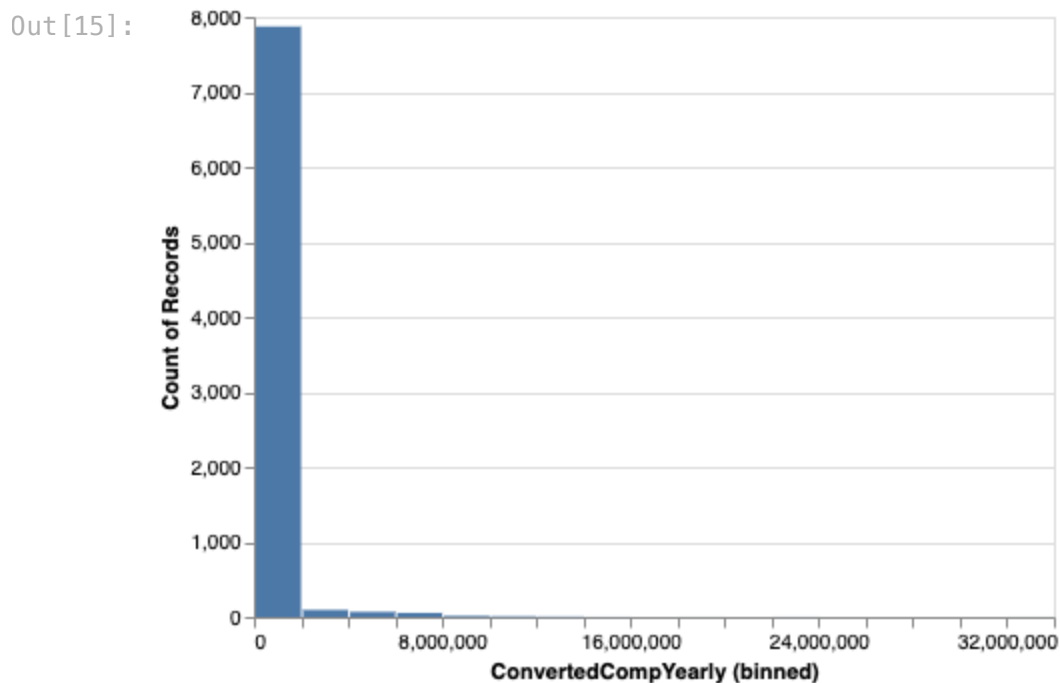






Taking the look at the distribution in our plots, all these features can be used in our model by one hot encoding them. There are a lot of missing values, we will try various imputation techniques to deal with this issue. For class imbalance issues, we will try to club classes having very few counts or make use undersampling/oversampling techniques

In [15]: `converted_comp_hist`




This shows that the salary data is highly skewed therefore it is better if we take median as a measure since it is better against the outliers

## 2D Histogram for YearsCode and ConvertedCompYearly

There appear to be very less outliers in the above chart and also most of the data lies between 6-12 years


```
In [16]: years_code
```



## 2D Histogram for YearsCodePro and ConvertedCompYearly

There appear to be very less outliers in the above chart and also most of the data lies between 2-8 years


```
In [17]: years_code
```



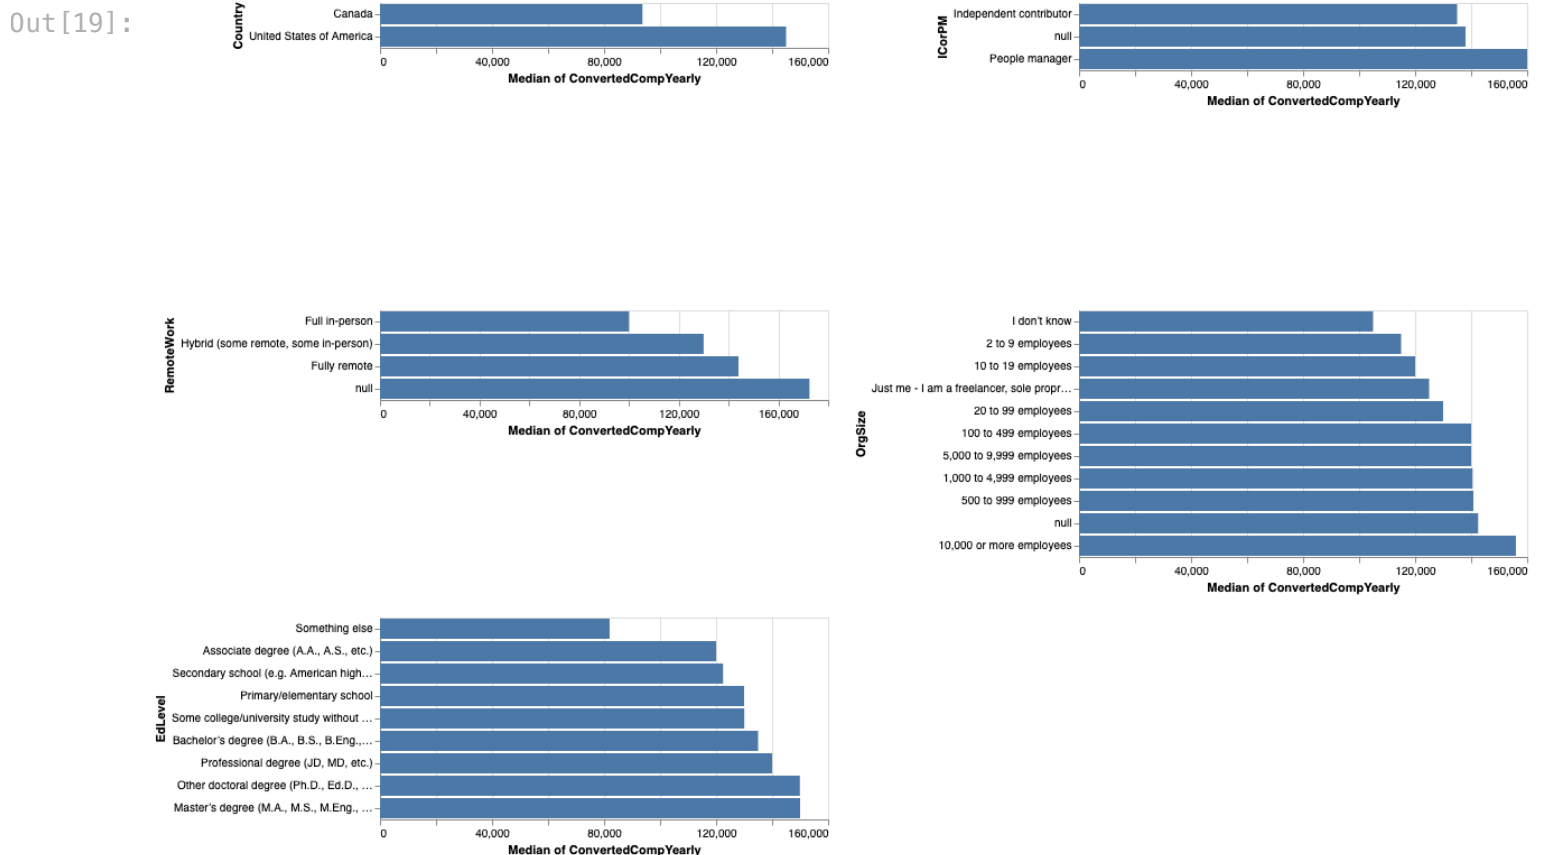
## 2D Histogram for WorkExp and ConvertedCompYearly

There appear to be barely any outliers in the above chart and also most of the data lies between 2-6 years

```
In [18]: work_exp
```



```
In [19]: multi_bar_plot
```



1. The bar chart for Country and ConvertedCompYearly shows that median salary for US is about 130K and 95K for Canada

2. The bar chart for ICorPM and ConvertedCompYearly shows that highest median salary are for people manager
3. The bar chart for RemoteWork and ConvertedCompYearly shows that median salary for is highest for those working fully remote
4. The bar chart for Orgsize and ConvertedCompYearly shows that median salary is highest for companies having 10,000 employees or more
5. The bar chart for Edlevel and ConvertedCompYearly shows that highest median salary are having a masters degree

Overall here there is less class imbalance for these features

In [20]: `corr_table`

Out[20]:

Table 2. Correlation Plot between numeric data

	<b>YearsCode</b>	<b>YearsCodePro</b>	<b>WorkExp</b>	<b>ConvertedCompYearly</b>
<b>YearsCode</b>	1.000000	0.915677	0.844152	0.018527
<b>YearsCodePro</b>	0.915677	1.000000	0.902687	0.019431
<b>WorkExp</b>	0.844152	0.902687	1.000000	0.007428
<b>ConvertedCompYearly</b>	0.018527	0.019431	0.007428	1.000000

It is seen that there is very high correlation between the three numeric features so they can be used to predict the ConvertedCompYearly

## References

1. Analyst-2 (analyst-2.ai) / Inspirient GmbH (inspirient.com) (2021). 'Salary and more-Data Scientist, Analyst, Engineer' analyzed by Analyst-2 [Dataset]. <https://info.stackoverflowsolutions.com/rs/719-EMH-566/images/stack-overflow-developer-survey-2022.zip>
2. <https://www.kaggle.com/datasets/phuchuynguyen/salary-and-moredata-scientist-analyst-engineer>